

*Lab Report on*

# CSE 452 - Machine Learning Lab



## **Submitted To**

**Md. Ferdous**

Lecturer

Dept. of CSE, BSMRSTU, Gopalganj  
E-mail: [md.ferdous@bsmrstu.edu.bd](mailto:md.ferdous@bsmrstu.edu.bd)

## **Submitted By**

**Israt Jahan Reshma**

ID:18CSE241

4<sup>th</sup> Year, 2<sup>nd</sup> Semester

Session: 2018-2019

Dept. of CSE, BSMRSTU

E-mail: [israt.18cse241@bsmrstu.edu.bd](mailto:israt.18cse241@bsmrstu.edu.bd)

## **Department of Computer Science & Engineering**

Bangabandhu Sheikh Mujibur Rahman Science & Technology University, Gopalganj – 8100

*10 November 2024*

## **ML Algorithms and Used dataset overview**

1. **Data Analysis: IT Student Mental Survey Dataset (Kaggle)**
  - a. **Algorithm:** Data Analysis
  - b. **Purpose:** Analyze mental health challenges (stress, anxiety) faced by IT students.
2. **Linear Regression: Student Performance Factors Dataset (Kaggle)**
  - a. **Algorithm:** Linear Regression
  - b. **Purpose:** Predict student grades based on factors like demographics and study habits.
3. **Logistic Regression: Student Performance Dataset (Kaggle)**
  - a. **Algorithm:** Logistic Regression
  - b. **Purpose:** Predict pass/fail outcomes based on student characteristics.
4. **K Means Clustering: Daily Delhi Climate Dataset (Kaggle)**
  - a. **Algorithm:** K Means Clustering
  - b. **Purpose:** Categorize days into similar weather patterns (temperature, humidity, air quality).
5. **SVM: Hepatitis.csv Dataset (Kaggle)**
  - a. **Algorithm:** Support Vector Machine (SVM)
  - b. **Purpose:** Classify patients as having hepatitis or not based on medical data.
6. **ANN: Keras MNIST Dataset**
  - a. **Algorithm:** Artificial Neural Network (ANN)
  - b. **Purpose:** Classify handwritten digits (0-9) using image data.
7. **CNN: Cat and Dog Image Classification Dataset**
  - a. **Algorithm:** Convolutional Neural Network (CNN)
  - b. **Purpose:** Classify images as cat or dog.
8. **LSTM: DailyDelhiClimateTrain.csv and Test Dataset (Kaggle)**
  - a. **Algorithm:** Long Short-Term Memory (LSTM)
  - b. **Purpose:** Predict future climate trends based on historical data.

# Data Analysis

The Data set is taken from Kaggle named as IT student Mental Servay.

```
In [46]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [47]: df = pd.read_csv('MentalHealthSurveyfoITstudents.csv');
df.head()
```

	gender	age	university	degree_level	degree_major	academic_year	cgpa	residential_status	can
0	Male	20	PU	Undergraduate	Data Science	2nd year	3.0-3.5	Off-Campus	
1	Male	20	UET	Postgraduate	Computer Science	3rd year	3.0-3.5	Off-Campus	
2	Male	20	FAST	Undergraduate	Computer Science	3rd year	2.5-3.0	Off-Campus	
3	Male	20	UET	Undergraduate	Computer Science	3rd year	2.5-3.0	On-Campus	
4	Female	20	UET	Undergraduate	Computer Science	3rd year	3.0-3.5	Off-Campus	

5 rows × 21 columns

```
In [48]: print("Shape : ", df.shape)
```

Shape : (87, 21)

```
In [49]: df.duplicated().sum()
df.isna().sum().sum()
```

```
Out[49]: 0
```

```
In [50]: df.describe().T
```

Out[50]:

		count	mean	std	min	25%	50%	75%	max
	<b>age</b>	87.0	19.942529	1.623636	17.0	19.0	20.0	21.0	26.0
	<b>study_satisfaction</b>	87.0	3.931034	1.043174	1.0	3.0	4.0	5.0	5.0
	<b>academic_workload</b>	87.0	3.885057	0.854880	2.0	3.0	4.0	4.5	5.0
	<b>academic_pressure</b>	87.0	3.781609	1.125035	1.0	3.0	4.0	5.0	5.0
	<b>financial_concerns</b>	87.0	3.390805	1.400634	1.0	2.5	3.0	5.0	5.0
	<b>social_relationships</b>	87.0	2.781609	1.175578	1.0	2.0	3.0	4.0	5.0
	<b>depression</b>	87.0	3.218391	1.367609	1.0	2.0	3.0	4.0	5.0
	<b>anxiety</b>	87.0	3.218391	1.297809	1.0	2.0	3.0	4.0	5.0
	<b>isolation</b>	87.0	3.241379	1.405682	1.0	2.0	3.0	4.5	5.0
	<b>future_insecurity</b>	87.0	3.011494	1.385089	1.0	2.0	3.0	4.0	5.0

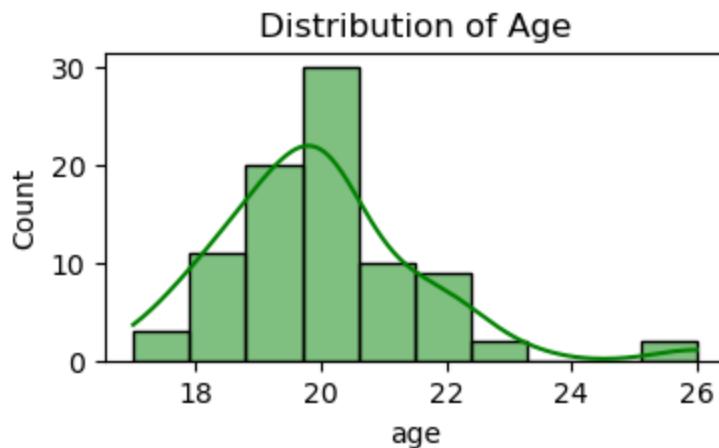
In [51]: `df['gender'].value_counts().T`

Out[51]: gender

Male 63

Female 24

Name: count, dtype: int64

In [52]: `plt.figure(figsize=(4,2))  
sns.histplot(df['age'], kde=True, color='green', edgecolor="black")  
plt.title("Distribution of Age")  
plt.xlabel("age")  
plt.show()`In [53]: `df['degree_level'].value_counts().T`

Out[53]: degree\_level

Undergraduate 85

Postgraduate 2

Name: count, dtype: int64

In [54]: `df['degree_major'].value_counts().T`

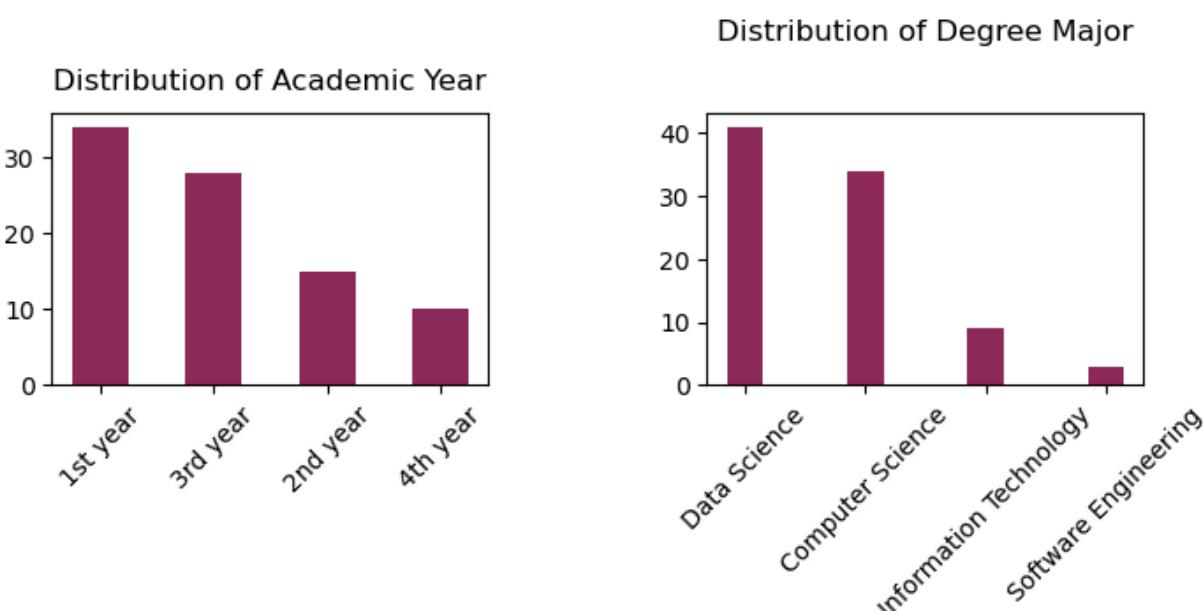
```
Out[54]: degree_major
Data Science          41
Computer Science      34
Information Technology 9
Software Engineering   3
Name: count, dtype: int64
```

```
In [56]: fig, axs = plt.subplots(1, 2, figsize=(8, 2))

axs[0].bar(df['academic_year'].value_counts().index,
            df['academic_year'].value_counts().values,
            color="#8f285b", width=0.5)
axs[0].set_title("Distribution of Academic Year", pad=10)
axs[0].tick_params(axis='x', rotation=45)

axs[1].bar(df['degree_major'].value_counts().index,
            df['degree_major'].value_counts().values,
            color="#8f285b", width=0.3)
axs[1].set_title("Distribution of Degree Major", pad=30)
axs[1].tick_params(axis='x', rotation=45)

plt.subplots_adjust(wspace=0.5)
plt.show()
```



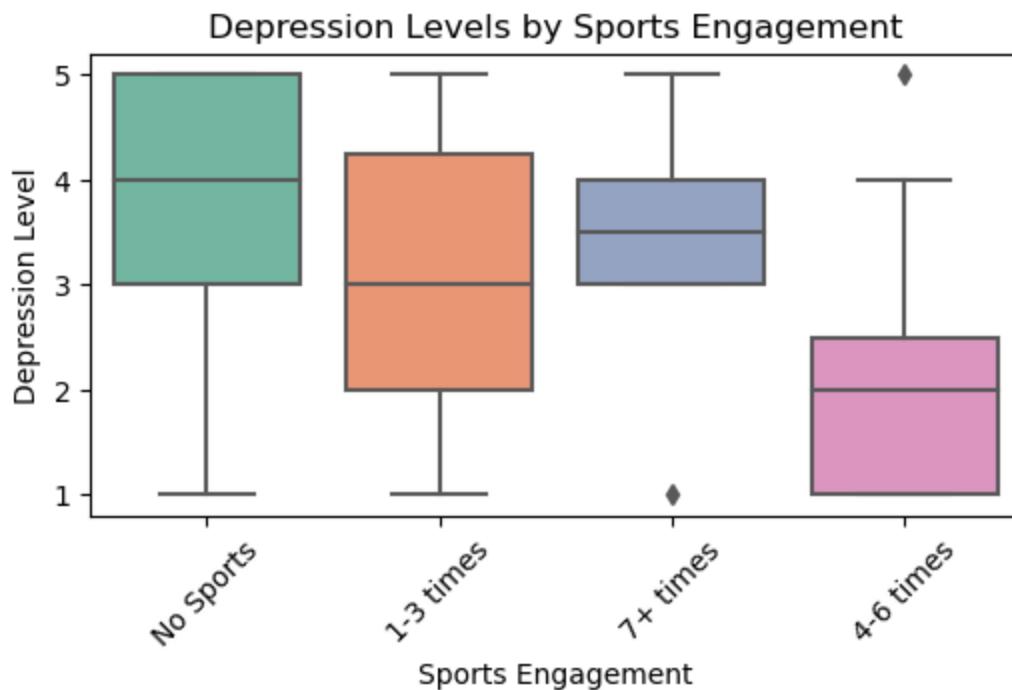
```
In [57]: df.cgpa.value_counts().T
```

```
Out[57]: cgpa
3.0-3.5    27
3.5-4.0    27
2.5-3.0    22
0.0-0.0     5
2.0-2.5     4
1.5-2.0     2
Name: count, dtype: int64
```

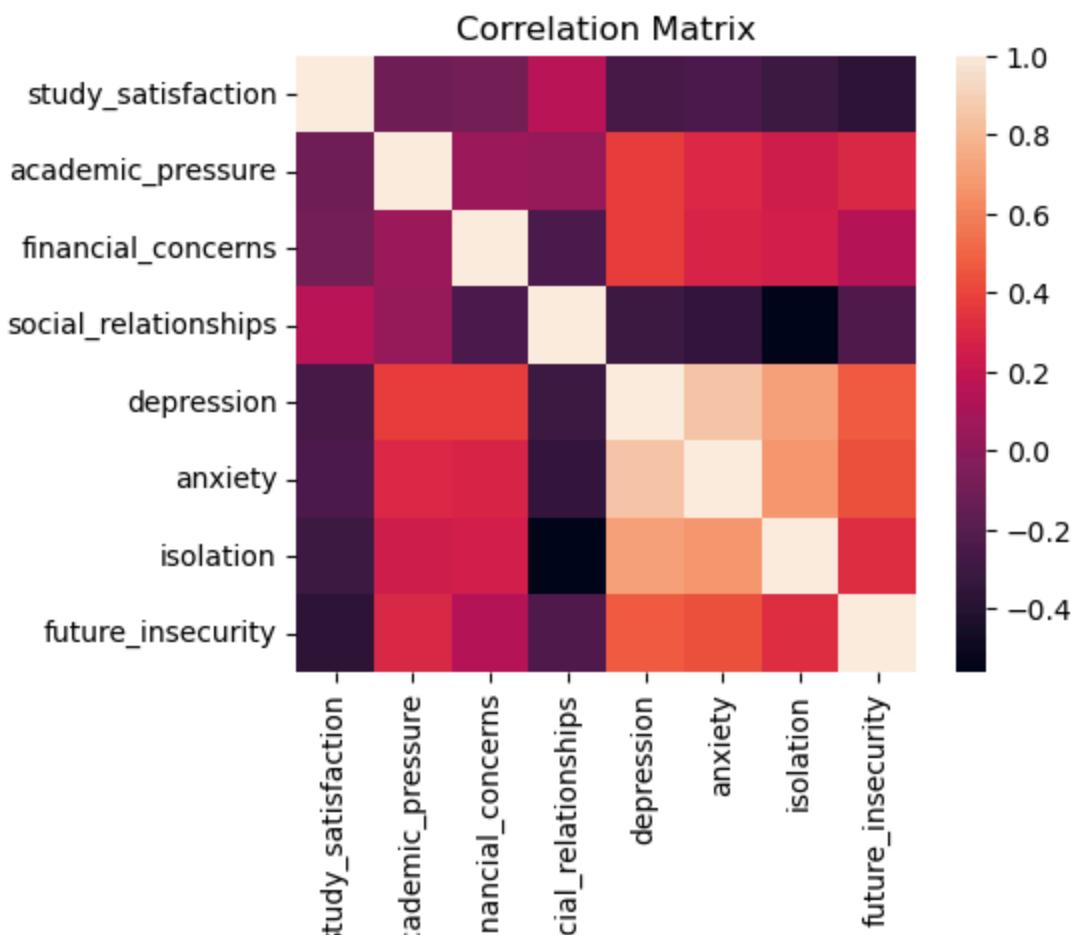
```
In [58]: df.residential_status.value_counts().T
```

```
Out[58]: residential_status  
Off-Campus      65  
On-Campus       22  
Name: count, dtype: int64
```

```
In [59]: plt.figure(figsize=(6, 3))  
sns.boxplot(x='sports_engagement', y='depression', data=df, palette="Set2")  
plt.title('Depression Levels by Sports Engagement')  
plt.xlabel('Sports Engagement')  
plt.ylabel('Depression Level')  
plt.xticks(rotation=45)  
plt.show()
```



```
In [60]: data_corr=df[['study_satisfaction','academic_pressure','financial_concerns','social_re  
#data_corr  
plt.figure(figsize=(5, 4))  
sns.heatmap(data_corr)  
plt.title('Correlation Matrix')  
plt.show()
```



## Linear Regression

Take Student Performance Factors Dataset from Kaggle

```
In [66]: df = pd.read_csv('StudentPerformanceFactors.csv')
df.head(2)
```

```
Out[66]:
```

	Hours_Studied	Attendance	Parental_Involvement	Access_to_Resources	Extracurricular_Activities	SI
0	23	84		Low	High	No
1	19	64		Low	Medium	No

```
In [67]: df=df.drop_duplicates()
#print(df.isna().sum())
df=df.dropna()
```

```
In [68]: df.describe()
```

Out[68]:	Hours_Studied	Attendance	Sleep_Hours	Previous_Scores	Tutoring_Sessions	Physical_Activity
<b>count</b>	6378.000000	6378.000000	6378.000000	6378.000000	6378.000000	6378.000000
<b>mean</b>	19.977109	80.020853	7.034964	75.066165	1.495296	2.972719
<b>std</b>	5.985460	11.550723	1.468033	14.400389	1.233984	1.028926
<b>min</b>	1.000000	60.000000	4.000000	50.000000	0.000000	0.000000
<b>25%</b>	16.000000	70.000000	6.000000	63.000000	1.000000	2.000000
<b>50%</b>	20.000000	80.000000	7.000000	75.000000	1.000000	3.000000
<b>75%</b>	24.000000	90.000000	8.000000	88.000000	2.000000	4.000000
<b>max</b>	44.000000	100.000000	10.000000	100.000000	8.000000	6.000000

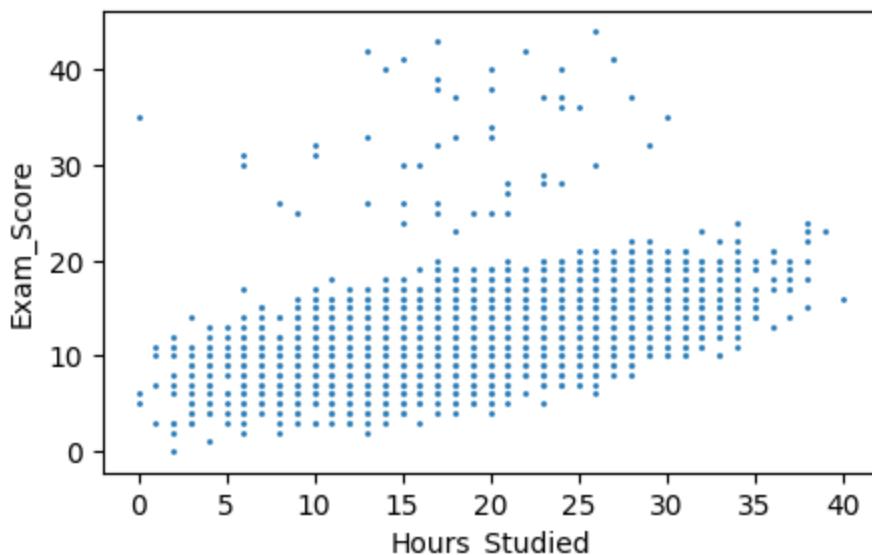
In [69]:

```
from sklearn.preprocessing import LabelEncoder
label_encoder = LabelEncoder()
for i in df.columns:
    df[i] = label_encoder.fit_transform(df[i])
df.head(2)
```

Out[69]:	Hours_Studied	Attendance	Parental_Involvement	Access_to_Resources	Extracurricular_Activities	SI
<b>0</b>	22	24		1	0	0
<b>1</b>	18	4		1	2	0

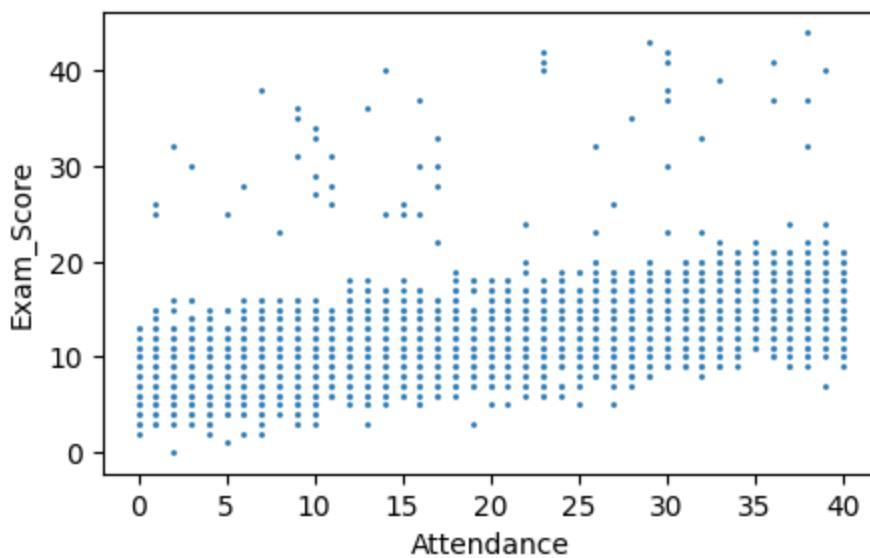
In [97]:

```
plt.figure(figsize=(5, 3)) # Adjust the figure size as needed
sns.scatterplot(data=df, x='Hours_Studied', y='Exam_Score', s=5)
plt.show()
```



In [98]:

```
plt.figure(figsize=(5, 3))
sns.scatterplot(data=df, x='Attendance', y='Exam_Score', s=5)
plt.show()
```



In [99]:

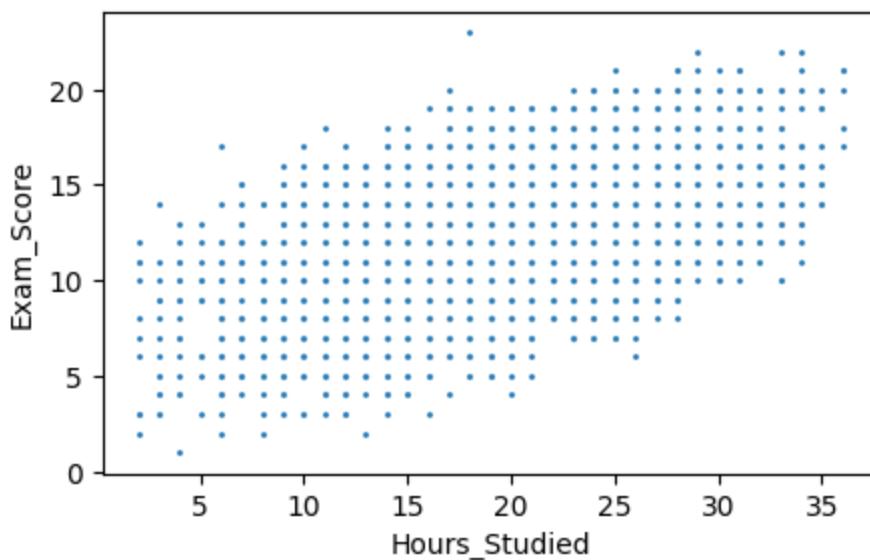
```
#removing outliers
mean = np.mean(df, axis=0)
std_dev = np.std(df, axis=0)
z_scores = (df - mean) / std_dev
outliers = np.abs(z_scores) > 3
df = df[~np.any(outliers, axis=1)]
df.head()
```

Out[99]:

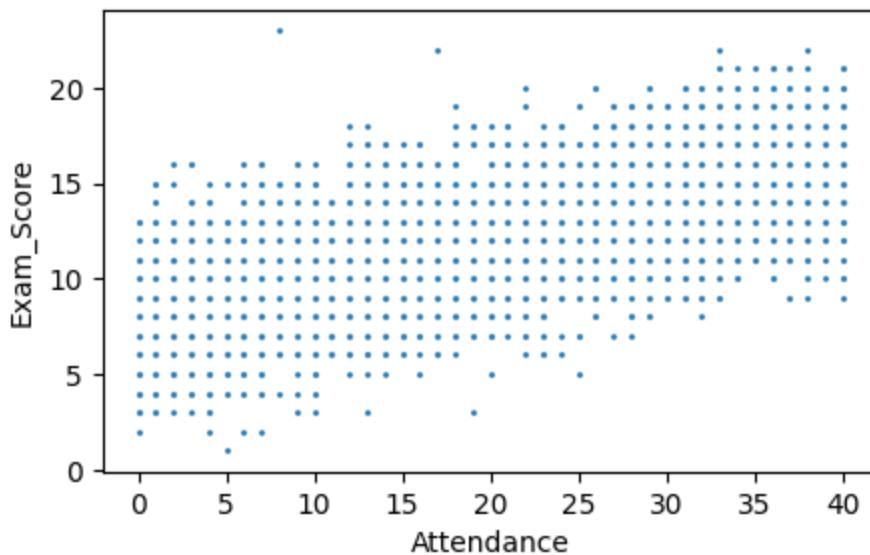
	Hours_Studied	Attendance	Parental_Involvement	Access_to_Resources	Extracurricular_Activities	SI
0	22	24	1	0	0	0
1	18	4	1	2	0	0
2	23	38	2	2	1	1
3	28	29	1	2	1	1
4	18	32	2	2	1	1

In [100...]

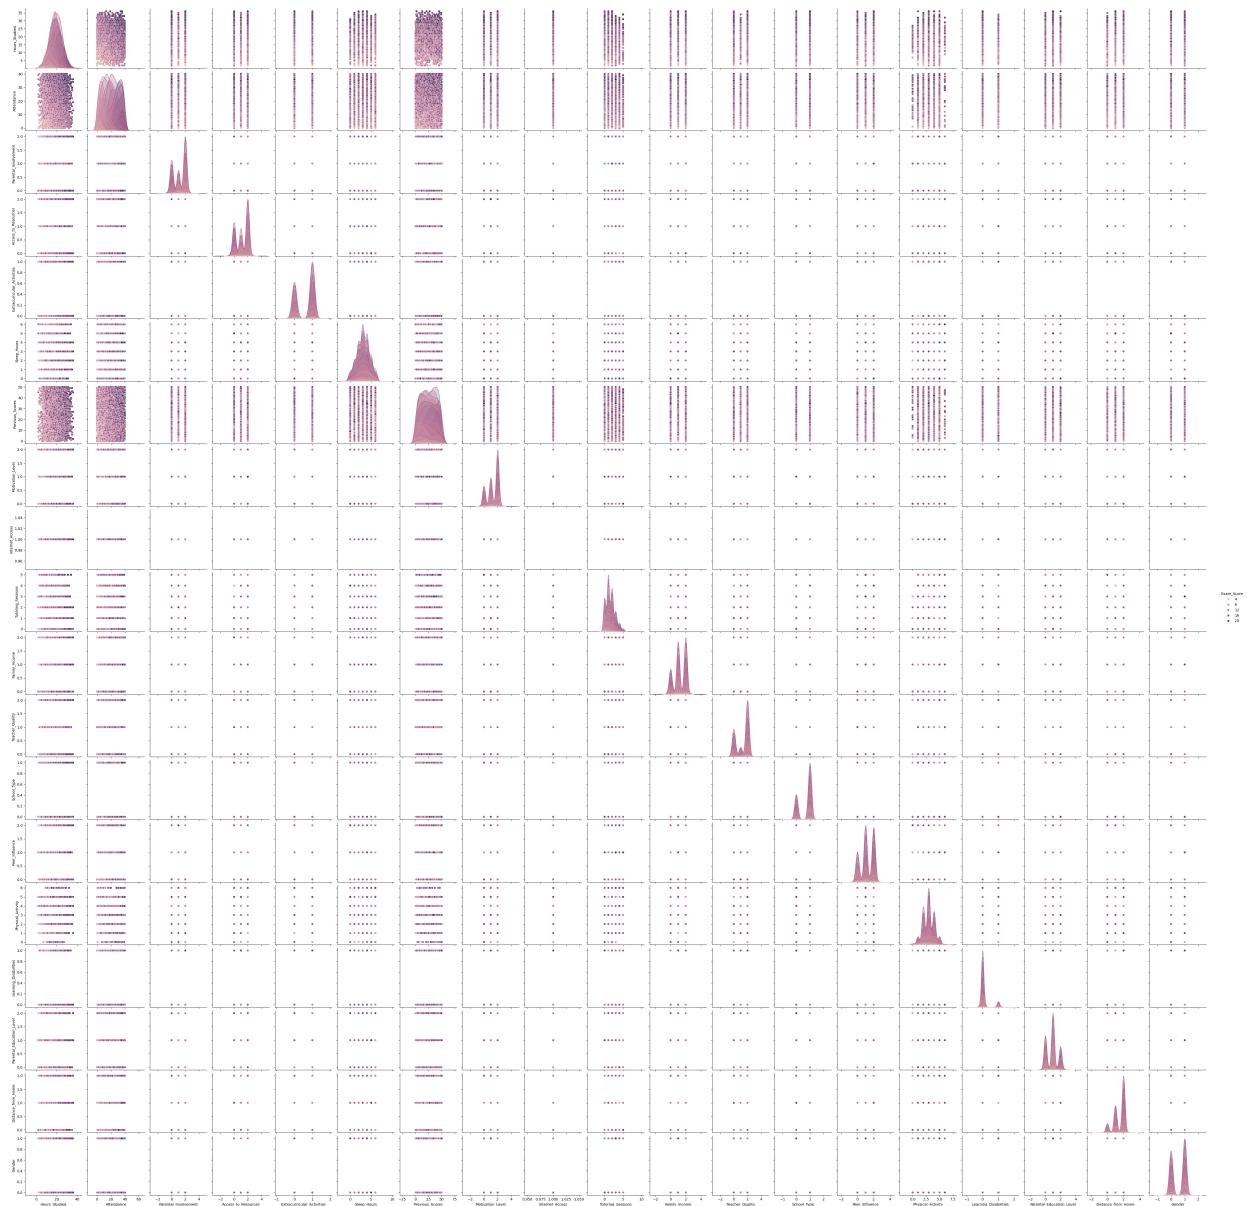
```
plt.figure(figsize=(5, 3))
sns.scatterplot(data=df, x='Hours_Studied', y='Exam_Score', s=5)
plt.show()
```



```
In [101...]:  
plt.figure(figsize=(5, 3))  
sns.scatterplot(data=df, x='Attendance', y='Exam_Score', s=5)  
plt.show()
```

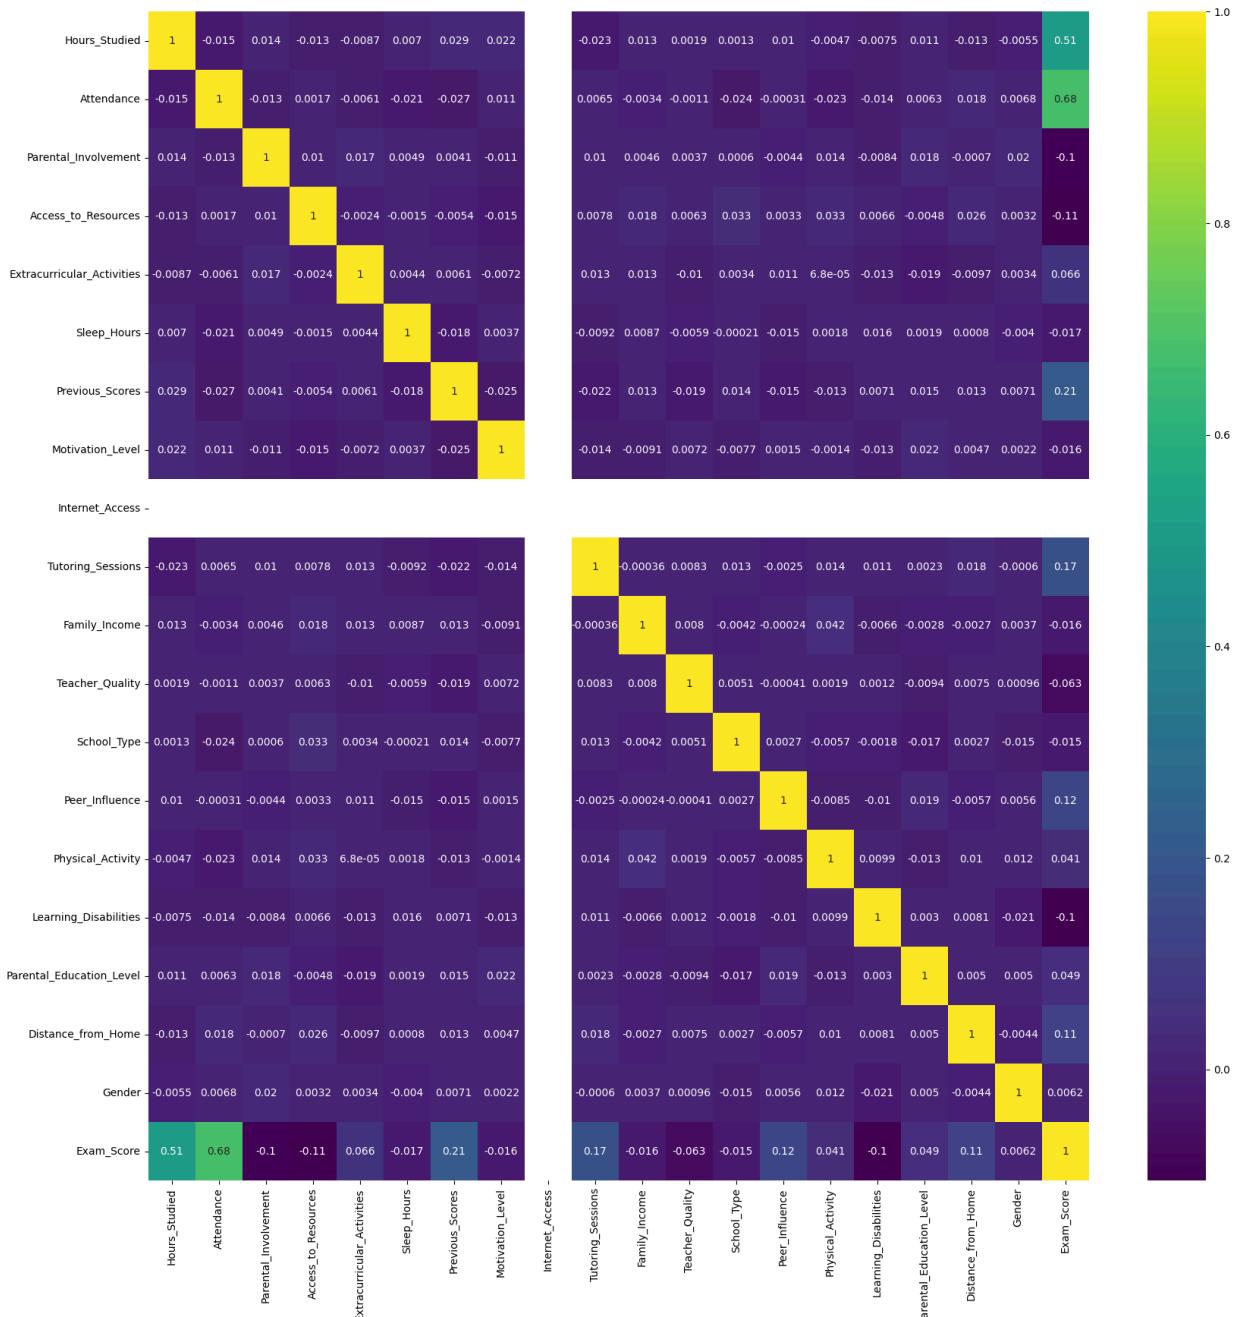


```
In [102...]:  
sns.pairplot(df, hue="Exam_Score")  
plt.show()
```



In [103...]

```
Correlation_Matrix=df.corr()
plt.figure(figsize=(20,20))
sns.heatmap(Correlation_Matrix, annot=True, cmap='viridis')
plt.show()
```



```
In [104]: X=df[['Hours_Studied', 'Attendance']]
y=df['Exam_Score']
```

```
In [105]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=42)
```

```
In [106]: from sklearn.linear_model import LinearRegression
reg=LinearRegression()
```

```
In [107]: reg.fit(x_train,y_train)
```

```
Out[107]: ▾ LinearRegression ⓘ ?
```

```
LinearRegression()
```

```
In [108]: y_pred=reg.predict(x_test)

In [109]: from sklearn.metrics import r2_score
r2_score(y_test,y_pred)

Out[109]: 0.7263622686013315
```

## Logistic Regression

take student performance data set from kaggle

```
In [110]: df = pd.read_csv('student_performance.csv')

In [111]: df.head()

Out[111]:
```

	Marital status	Application mode	Application order	Course	Daytime/evening attendance	Previous qualification	Nacionality	Mother's qualification
0	1	8	5	2	1	1	1	13
1	1	6	1	11	1	1	1	1
2	1	1	5	5	1	1	1	22
3	1	8	2	15	1	1	1	23
4	2	12	1	3	0	1	1	22

5 rows × 35 columns



```
In [112]: df.isna().sum().sum()

Out[112]: 0
```

```
In [113]: df['Target'].value_counts()
```

```
Out[113]:
```

Target	count
Graduate	2209
Dropout	1421
Enrolled	794

**dtype:** int64

```
In [114]: df['Target']=df['Target'].apply(lambda x: 1 if x=='Graduate' else 0)
df['Target'].value_counts()
```

Out[114]: count

**Target**

0	2215
1	2209

**dtype:** int64

```
In [115]: X=df.drop('Target', axis=1)
y=df['Target']
```

```
In [116]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
X_train, X_test, Y_train, Y_test = train_test_split(X, y, test_size=0.2, random_state=42)
logmodel=LogisticRegression(max_iter=5000)
logmodel.fit(X_train, Y_train)
```

Out[116]:

**LogisticRegression** ⓘ ?  
**LogisticRegression(max\_iter=5000)**

```
In [117]: predictions=logmodel.predict(X_test)
```

```
In [118]: from sklearn.metrics import classification_report
print(classification_report(Y_test, predictions))
```

	precision	recall	f1-score	support
0	0.89	0.79	0.84	467
1	0.79	0.89	0.84	418
accuracy			0.84	885
macro avg	0.84	0.84	0.84	885
weighted avg	0.84	0.84	0.84	885

## K means Clustering

Data set is taken from Kaggle : name: Daily Delhi Climate dataset

```
In [101]: df=pd.read_csv("DailyDelhiClimateTrain.csv")
df.head()
```

	date	meantemp	humidity	wind_speed	meanpressure
0	2013-01-01	10.000000	84.500000	0.000000	1015.666667
1	2013-01-02	7.400000	92.000000	2.980000	1017.800000
2	2013-01-03	7.166667	87.000000	4.633333	1018.666667
3	2013-01-04	8.666667	71.333333	1.233333	1017.166667
4	2013-01-05	6.000000	86.833333	3.700000	1016.500000

In [102...]: `df.shape`

Out[102]: (1462, 5)

In [103...]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1462 entries, 0 to 1461
Data columns (total 5 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   date        1462 non-null   object 
 1   meantemp    1462 non-null   float64
 2   humidity    1462 non-null   float64
 3   wind_speed  1462 non-null   float64
 4   meanpressure 1462 non-null   float64
dtypes: float64(4), object(1)
memory usage: 57.2+ KB
```

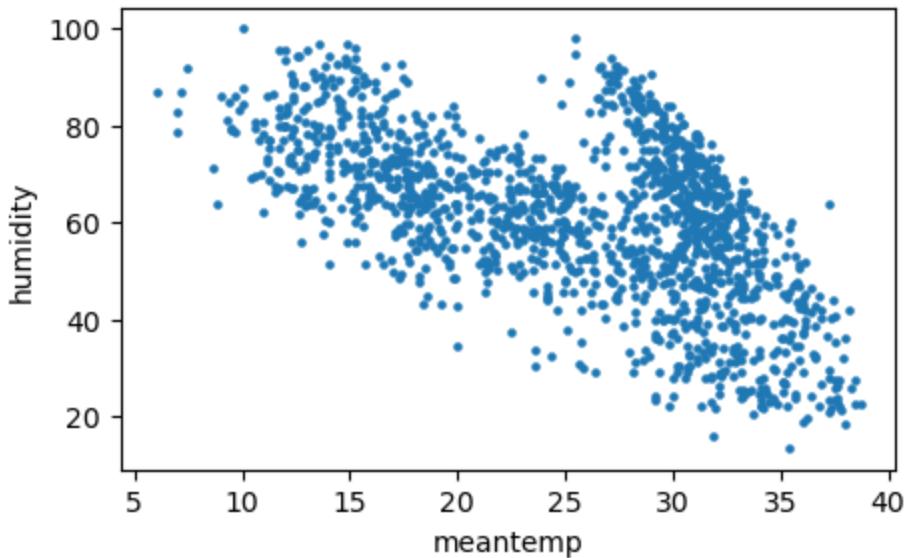
In [104...]: `df.describe()`

	meantemp	humidity	wind_speed	meanpressure
count	1462.000000	1462.000000	1462.000000	1462.000000
mean	25.495521	60.771702	6.802209	1011.104548
std	7.348103	16.769652	4.561602	180.231668
min	6.000000	13.428571	0.000000	-3.041667
25%	18.857143	50.375000	3.475000	1001.580357
50%	27.714286	62.625000	6.221667	1008.563492
75%	31.305804	72.218750	9.238235	1014.944901
max	38.714286	100.000000	42.220000	7679.333333

In [113...]: `df.plot.scatter(x='meantemp', y='humidity', title='Mean Temp vs Humidity', s=5, figsize=(10, 6))`

Out[113]: <Figure>

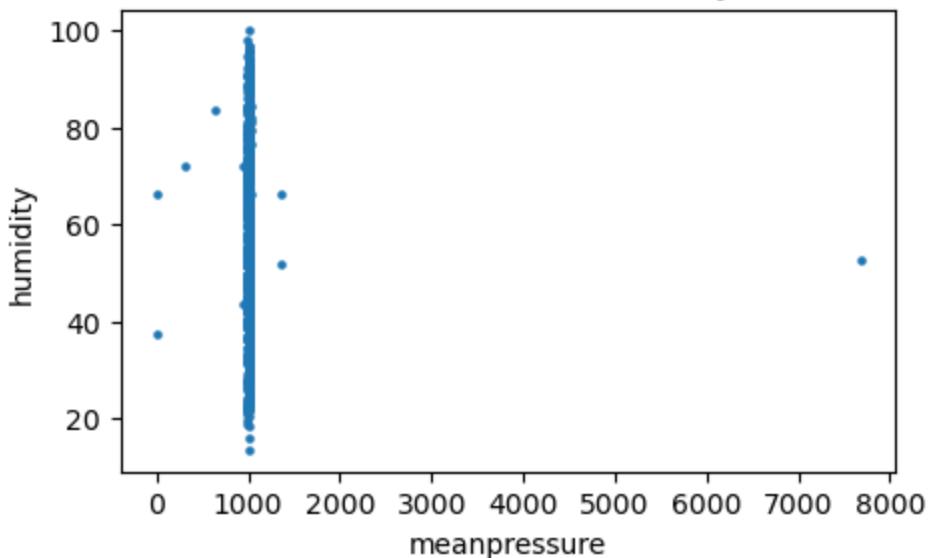
### Mean Temp vs Humidity



```
In [112]: df.plot.scatter(x='meanpressure', y='humidity', title='Mean Pressure vs Humidity', s=5)
```

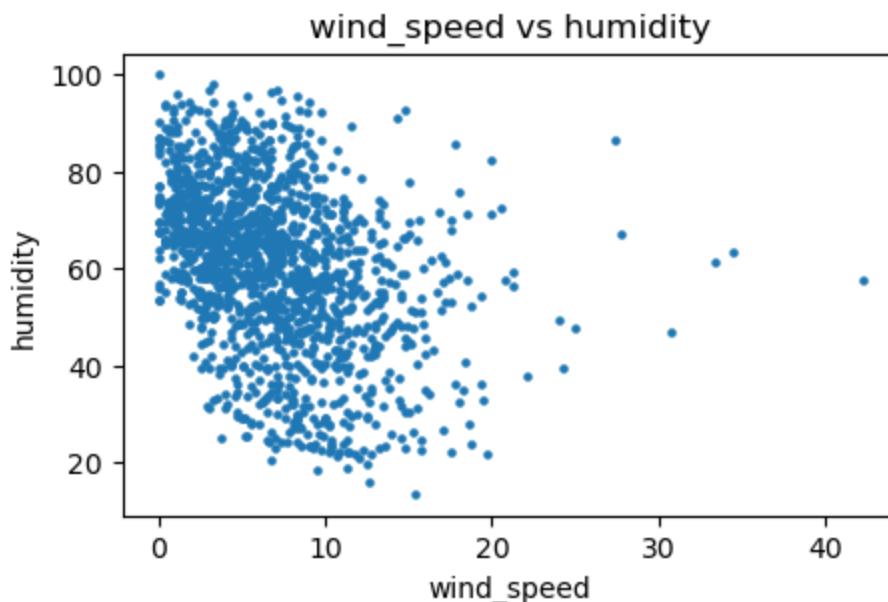
```
Out[112]: <Axes: title={'center': 'Mean Pressure vs Humidity'}, xlabel='meanpressure', ylabel='humidity'>
```

### Mean Pressure vs Humidity



```
In [114]: df.plot(x='wind_speed',
                 y='humidity',
                 kind='scatter',
                 title='wind_speed vs humidity', s=5, figsize=(5, 3))
```

```
Out[114]: <Axes: title={'center': 'wind_speed vs humidity'}, xlabel='wind_speed', ylabel='humidity'>
```



In [115...]

```
from sklearn.cluster import KMeans
import plotly.express as px

x= df[['humidity','meanpressure','meantemp','wind_speed']]
km = KMeans(n_clusters=3, n_init=10)
df['cluster']= km.fit_predict(x)
df['cluster']=df['cluster'] \
    .astype('category')
```

In [118...]

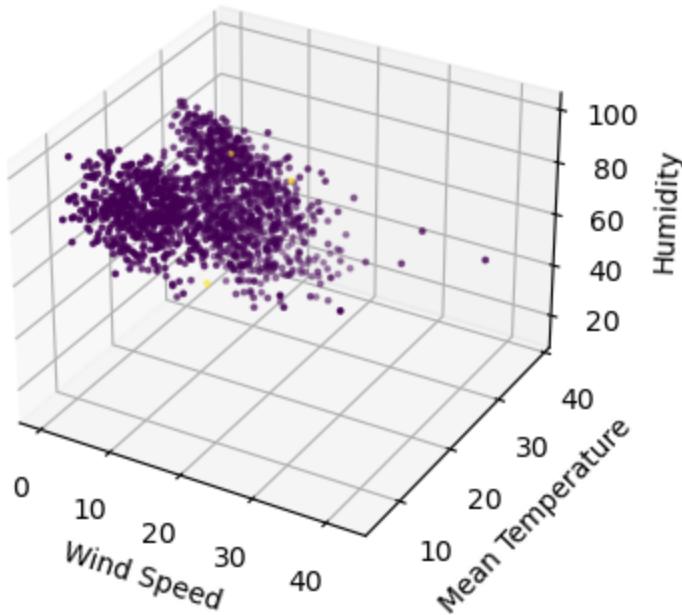
```
fig = plt.figure(figsize=(6, 4))
ax = fig.add_subplot(111, projection='3d')

# Plot the 3D scatter plot
ax.scatter(df['wind_speed'], df['meantemp'], df['humidity'], c=df['cluster'].astype('category'))

# Set the labels and title
ax.set_xlabel('Wind Speed')
ax.set_ylabel('Mean Temperature')
ax.set_zlabel('Humidity')
ax.set_title('3D Scatter Plot: Wind Speed vs Temperature vs Humidity')

# Show the plot
plt.show()
```

### 3D Scatter Plot: Wind Speed vs Temperature vs Humidity



```
In [117]: df.drop(columns=['date'], inplace=True)
```

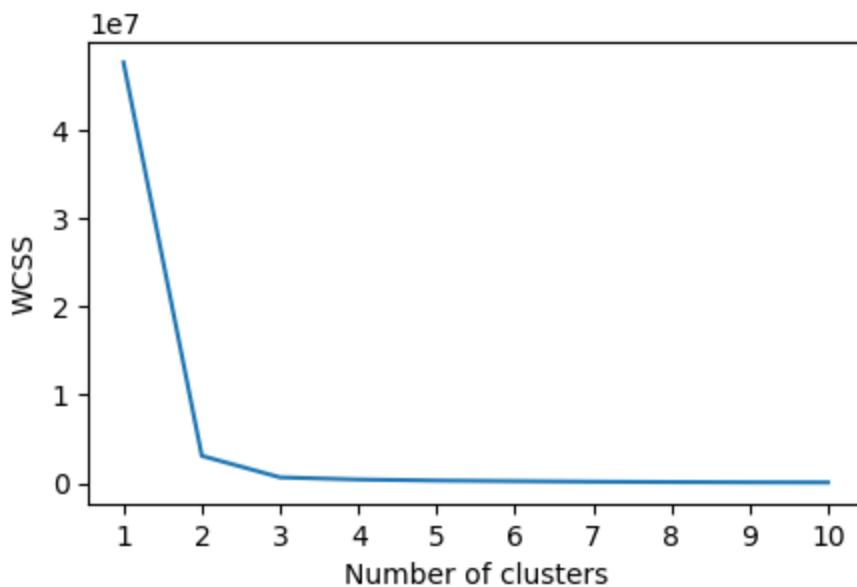
```
In [32]: x=df.drop(columns=['wind_speed'])
y=df['wind_speed']
```

```
In [35]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=42)
```

```
In [36]: wcss = []
for k in range(1, 11):
    kmeans = KMeans(n_clusters=k, init='k-means++', n_init=10)
    kmeans.fit(x_train)
    wcss.append(kmeans.inertia_)
wcss
```

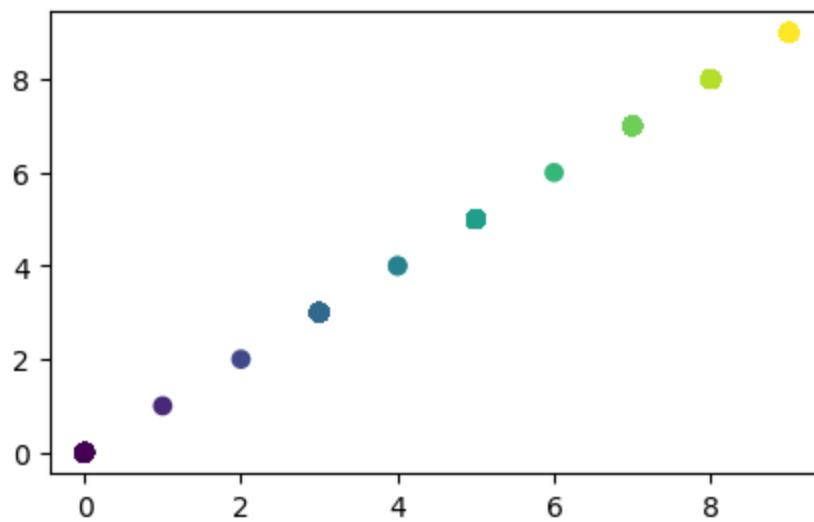
```
Out[36]: [47572716.82843734,
3078058.3838992873,
644083.7667092744,
400171.0454776609,
284528.8565589274,
216098.55890250346,
153542.54396055674,
109239.52061460483,
82544.78736091012,
70091.01770029642]
```

```
In [38]: plt.figure(figsize=(5, 3)) # Specify smaller width and height in inches
plt.plot(range(1, 11), wcss)
plt.xticks(range(1, 11))
plt.xlabel('Number of clusters')
plt.ylabel('WCSS') # Add ylabel if needed
plt.show()
```



```
In [39]: x_pred=kmeans.fit_predict(x_train)
```

```
In [40]: plt.figure(figsize=(5, 3)) # Specify the width and height in inches
plt.scatter(x_pred, x_pred, c=x_pred)
plt.show()
```



```
In [47]: from kneed import KneeLocator
k1= KneeLocator(range(1,11),wcss,curve='convex',direction='decreasing')
k1.elbow
```

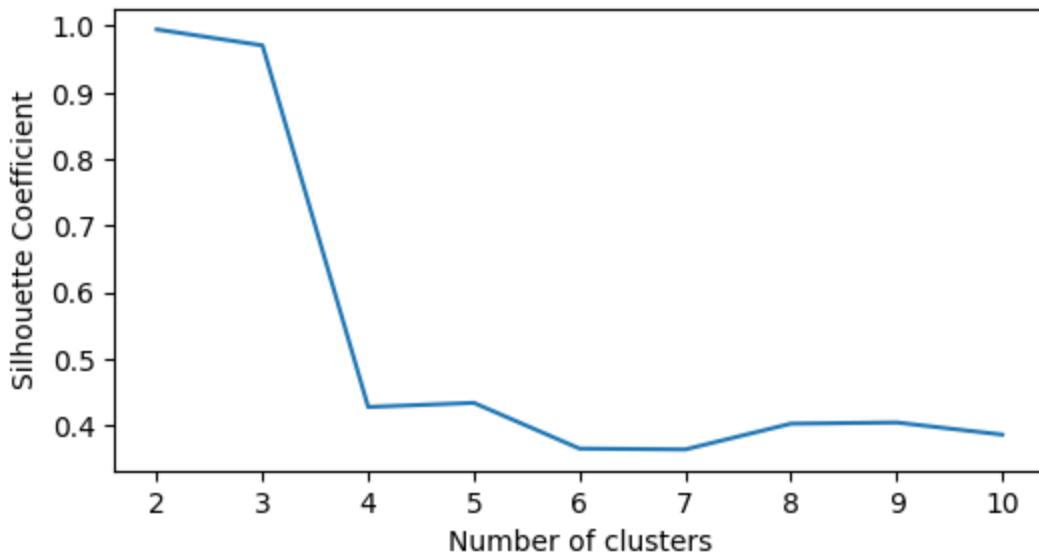
Out[47]: 2

```
In [48]: #performance metrics
from sklearn.metrics import silhouette_score
silhouette_coefficients=[]
for k in range(2,11):
    kmeans= KMeans(n_clusters=k,init="k-means++",n_init=10)
    kmeans.fit(x_train)
    score=silhouette_score(x_train,kmeans.labels_)
    silhouette_coefficients.append(score)
```

```
In [49]: silhouette_coefficients
```

```
Out[49]: [0.9947409709877785,
 0.9707440447233446,
 0.4274533839989495,
 0.4336048534773166,
 0.36487063247390333,
 0.3636351197961745,
 0.4023410404344692,
 0.404124154769077,
 0.3859009921674613]
```

```
In [50]: plt.figure(figsize=(6, 3)) # Adjust the width and height as needed
plt.plot(range(2, 11), silhouette_coefficients)
plt.xticks(range(2, 11))
plt.xlabel('Number of clusters')
plt.ylabel('Silhouette Coefficient')
plt.show()
```



## SVM

use "hepatitis.csv" from Kaggle

```
In [5]: ## Read "hepatitis.csv" using pandas
# target = 1: Die; 2: Live
data = pd.read_csv("hepatitis.csv", na_values="?")
data.head()
```

	class	age	sex	steroid	antivirals	fatigue	malaise	anorexia	liver_big	liver_firm	spleen_palable
0	2	30	2	1	2	2	2	2	1	2	2
1	2	50	1	1	2	1	2	2	1	2	2
2	2	78	1	2	2	1	2	2	2	2	2
3	2	34	1	2	2	2	2	2	2	2	2
4	2	34	1	2	2	2	2	2	2	2	2

In [6]: `data.shape`

Out[6]: `(142, 20)`

In [7]: `data.describe()`

	class	age	sex	steroid	antivirals	fatigue	malaise	anorexia
count	142.000000	142.000000	142.000000	142.000000	142.000000	142.000000	142.000000	142.000000
mean	1.816901	40.816901	1.105634	1.514085	1.838028	1.359155	1.619718	1.795775
std	0.388116	12.189182	0.308456	0.501571	0.369729	0.481451	0.487174	0.404561
min	1.000000	7.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000
25%	2.000000	32.000000	1.000000	1.000000	2.000000	1.000000	1.000000	2.000000
50%	2.000000	39.000000	1.000000	2.000000	2.000000	1.000000	2.000000	2.000000
75%	2.000000	50.000000	1.000000	2.000000	2.000000	2.000000	2.000000	2.000000
max	2.000000	78.000000	2.000000	2.000000	2.000000	2.000000	2.000000	2.000000

In [8]: `data['class'].value_counts()`

Out[8]: `count`

class
2 116
1 26

**dtype:** int64

In [9]: `cat_cols = data.columns[data.nunique() < 5]`

In [10]: `num_cols = data.columns[data.nunique() >= 5]`

In [11]: `num_cols = ["age", "bilirubin", "alk_phosphate", "albumin", "protime"]`  
`cat_cols = ['sex', 'steroid', 'antivirals', 'fatigue', 'malaise', 'anorexia', 'liver_b`

```
In [12]: data.isna().sum().sum()
```

```
Out[12]: 0
```

```
In [13]: X = data.drop(['class'], axis = 1)
y = data['class']
print(X.shape, y.shape)
```

```
(142, 19) (142,)
```

```
In [15]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state=42)
```

```
In [16]: X_train[cat_cols] = X_train[cat_cols].astype('int')

# Test
X_test[cat_cols] = X_test[cat_cols].astype('int')
```

```
In [17]: X_train = pd.get_dummies(X_train, columns=cat_cols, drop_first=True) # Train
X_test = pd.get_dummies(X_test, columns=cat_cols, drop_first=True) # Test
```

```
In [18]: X_train.columns
```

```
Out[18]: Index(['age', 'bilirubin', 'alk_phosphate', 'sgot', 'albumin', 'prottime',
       'sex_2', 'steroid_2', 'antivirals_2', 'fatigue_2', 'malaise_2',
       'anorexia_2', 'liver_big_2', 'liver_firm_2', 'spleen_palable_2',
       'spiders_2', 'ascites_2', 'varices_2', 'histology_2'],
      dtype='object')
```

```
In [19]: from sklearn.preprocessing import StandardScaler
```

```
In [21]: from sklearn.preprocessing import StandardScaler

# Ensure the columns in num_cols are of dtype float64 for compatibility with scaler transform
X_train.loc[:, num_cols] = X_train.loc[:, num_cols].astype('float64')
X_test.loc[:, num_cols] = X_test.loc[:, num_cols].astype('float64')

# Initialize and fit the scaler
scaler = StandardScaler()
scaler.fit(X_train.loc[:, num_cols])

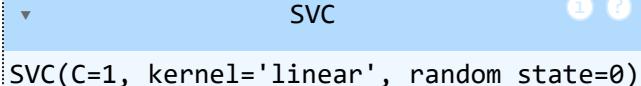
# Apply the transformation
X_train.loc[:, num_cols] = scaler.transform(X_train.loc[:, num_cols])
X_test.loc[:, num_cols] = scaler.transform(X_test.loc[:, num_cols])
```

## MODEL BUILDING - SVM

```
In [22]: from sklearn.svm import SVC
linear_svm = SVC(kernel='linear', C=1, random_state=0)
```

```
In [23]: # Train the classifier
linear_svm.fit(X=X_train, y=y_train)
```

Out[23]:



SVC(C=1, kernel='linear', random\_state=0)

In [24]:

```
## Predict
train_predictions = linear_svm.predict(X_train)
test_predictions = linear_svm.predict(X_test)

### Train data accuracy
from sklearn.metrics import accuracy_score,f1_score,confusion_matrix

print("TRAIN Conf Matrix : \n", confusion_matrix(y_train, train_predictions))
print("\nTRAIN DATA ACCURACY",accuracy_score(y_train,train_predictions))
print("\nTrain data f1-score for class '1'",f1_score(y_train,train_predictions, pos_label=1))
print("\nTrain data f1-score for class '2'",f1_score(y_train,train_predictions, pos_label=0))

### Test data accuracy
print("\n\n-----\n\n")

print("TEST Conf Matrix : \n", confusion_matrix(y_test, test_predictions))
print("\nTEST DATA ACCURACY",accuracy_score(y_test,test_predictions))
print("\nTest data f1-score for class '1'",f1_score(y_test,test_predictions, pos_label=1))
print("\nTest data f1-score for class '2'",f1_score(y_test,test_predictions, pos_label=0))
```

TRAIN Conf Matrix :

```
[[13  7]
 [ 3 90]]
```

TRAIN DATA ACCURACY 0.911504424778761

Train data f1-score for class '1' 0.7222222222222222

Train data f1-score for class '2' 0.9473684210526315

-----

TEST Conf Matrix :

```
[[ 1  5]
 [ 2 21]]
```

TEST DATA ACCURACY 0.7586206896551724

Test data f1-score for class '1' 0.2222222222222222

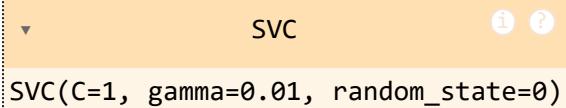
Test data f1-score for class '2' 0.8571428571428571

Kernel Trick:

In [25]:

```
svc = SVC(kernel='rbf', random_state=0, gamma=0.01, C=1)
svc
```

Out[25]:



SVC(C=1, gamma=0.01, random\_state=0)

```
In [26]: svc.fit(X=X_train, y= y_train)
```

Out[26]:

```
SVC(C=1, gamma=0.01, random_state=0)
```

```
In [27]: ## Predicttrain_predictions = svc.predict(X_train)
test_predictions = svc.predict(X_test)
```

```
### Train data accuracy
print("TRAIN Conf Matrix : \n", confusion_matrix(y_train, train_predictions))
print("\nTRAIN DATA ACCURACY",accuracy_score(y_train,train_predictions))
print("\nTrain data f1-score for class '1'",f1_score(y_train,train_predictions, pos_label=1))
print("\nTrain data f1-score for class '2'",f1_score(y_train,train_predictions, pos_label=0))
```

```
### Test data accuracy
print("\n\n-----\n\n")

print("TEST Conf Matrix : \n", confusion_matrix(y_test, test_predictions))
print("\nTEST DATA ACCURACY",accuracy_score(y_test,test_predictions))
print("\nTest data f1-score for class '1'",f1_score(y_test,test_predictions, pos_label=1))
print("\nTest data f1-score for class '2'",f1_score(y_test,test_predictions, pos_label=0))
```

TRAIN Conf Matrix :

```
[[13  7]
 [ 3 90]]
```

TRAIN DATA ACCURACY 0.911504424778761

Train data f1-score for class '1' 0.7222222222222222

Train data f1-score for class '2' 0.9473684210526315

-----

TEST Conf Matrix :

```
[[ 0  6]
 [ 0 23]]
```

TEST DATA ACCURACY 0.7931034482758621

Test data f1-score for class '1' 0.0

Test data f1-score for class '2' 0.8846153846153846

## ANN

Used keras mnist dataset

```
In [70]: import tensorflow
from tensorflow import keras
from tensorflow.keras import Sequential
from tensorflow.keras.layers import Dense, Flatten
```

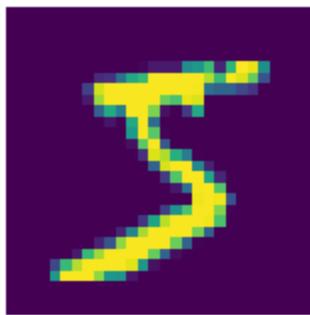
```
In [71]: (X_train,y_train),(X_test,y_test)=keras.datasets.mnist.load_data()
```

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz  
11490434/11490434 ━━━━━━━━ 0s 0us/step

```
In [72]: X_train.shape
```

```
Out[72]: (60000, 28, 28)
```

```
In [74]: plt.figure(figsize=(2, 2))
plt.imshow(X_train[0])
plt.axis('off')
plt.show()
```



```
In [75]: X_train=X_train/255
X_test=X_test/255
```

```
In [76]: #print(X_train[0])
```

```
In [77]: model = Sequential()
model.add(Flatten(input_shape=(28,28)))
model.add(Dense(128,activation='relu'))
model.add(Dense(32,activation='relu'))
model.add(Dense(128,activation='softmax'))
```

/usr/local/lib/python3.10/dist-packages/keras/src/layers/reshaping/flatten.py:37: UserWarning:

Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
In [78]: model.summary()
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Params
flatten (Flatten)	(None, 784)	16
dense (Dense)	(None, 128)	16
dense_1 (Dense)	(None, 32)	16
dense_2 (Dense)	(None, 128)	16

Total params: 108,832 (425.12 KB)

Trainable params: 108,832 (425.12 KB)

```
In [79]: model.compile(loss='sparse_categorical_crossentropy', optimizer='Adam', metrics=['accuracy'])
```

```
In [81]: history = model.fit(X_train, y_train, epochs=25, validation_split=0.2, verbose=0)
```

```
# Display the last 5 epochs' results
for epoch in range(-5, 0):
    print(f"Epoch {epoch + 25}: accuracy={history.history['accuracy'][epoch]:.4f}, "
          f"loss={history.history['loss'][epoch]:.4f}, "
          f"val_accuracy={history.history['val_accuracy'][epoch]:.4f}, "
          f"val_loss={history.history['val_loss'][epoch]:.4f}")
```

```
Epoch 20: accuracy=0.9976, loss=0.0082, val_accuracy=0.9743, val_loss=0.2065
Epoch 21: accuracy=0.9993, loss=0.0024, val_accuracy=0.9765, val_loss=0.2006
Epoch 22: accuracy=0.9980, loss=0.0071, val_accuracy=0.9734, val_loss=0.2172
Epoch 23: accuracy=0.9981, loss=0.0057, val_accuracy=0.9743, val_loss=0.2077
Epoch 24: accuracy=0.9983, loss=0.0055, val_accuracy=0.9750, val_loss=0.1973
```

```
In [82]: y_prob = model.predict(X_test)
```

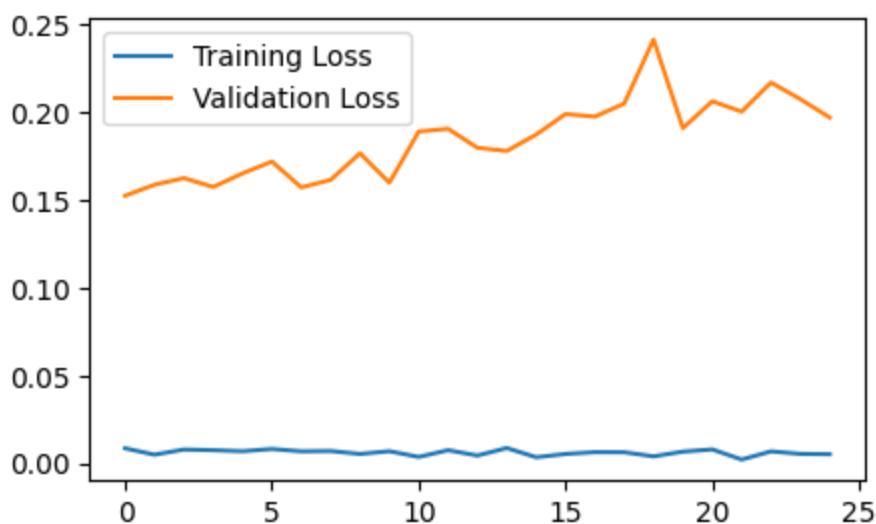
```
313/313 ━━━━━━━━ 1s 2ms/step
```

```
In [83]: y_pred = y_prob.argmax(axis=1)
```

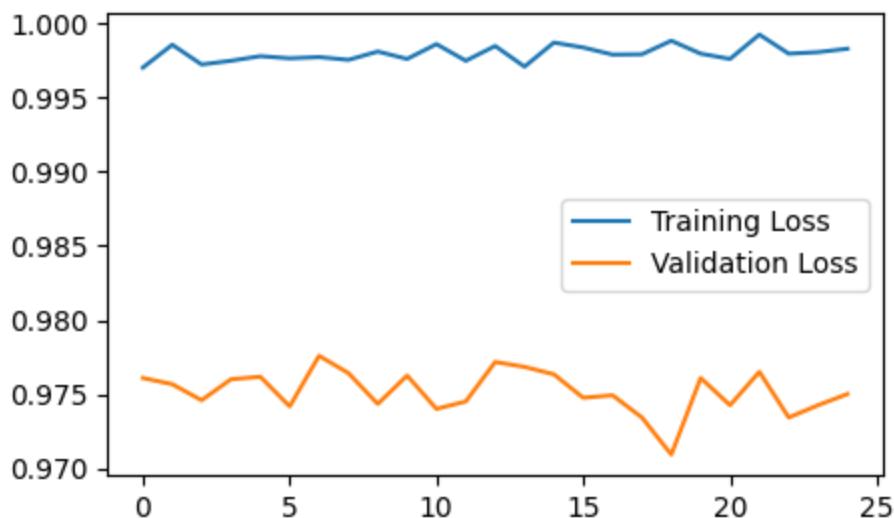
```
In [84]: from sklearn.metrics import accuracy_score
accuracy_score(y_test, y_pred)
```

```
Out[84]: 0.9766
```

```
In [90]: plt.figure(figsize=(5, 3)) # Set a smaller figure size
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.legend()
plt.show()
```



```
In [89]: plt.figure(figsize=(5, 3)) # Set a smaller figure size
plt.plot(history.history['accuracy'], label='Training Loss')
plt.plot(history.history['val_accuracy'], label='Validation Loss')
plt.legend()
plt.show()
```



```
In [92]: model.predict(X_test[1].reshape(1,28,28)).argmax(axis=1)
```

1/1 ————— 0s 24ms/step

Out[92]:

## CNN

Used cat dog image dataset for classifying cat and dog

```
In [2]: !mkdir -p ~/.kaggle
!cp kaggle.json ~/.kaggle/
```

cp: cannot stat 'kaggle.json': No such file or directory

```
In [3]: !kaggle datasets download -d salader/dogs-vs-cats
```

Dataset URL: <https://www.kaggle.com/datasets/salader/dogs-vs-cats>  
 License(s): unknown  
 Downloading dogs-vs-cats.zip to /content  
 100% 1.06G/1.06G [00:17<00:00, 64.5MB/s]  
 100% 1.06G/1.06G [00:17<00:00, 66.6MB/s]

```
In [4]: import zipfile
zip_ref=zipfile.ZipFile('/content/dogs-vs-cats.zip','r')
zip_ref.extractall('/content')
zip_ref.close()
```

```
In [5]: import tensorflow as tf
from tensorflow import keras
from keras import Sequential
from keras.layers import Dense, Flatten, Conv2D, MaxPooling2D, BatchNormalization, Dropout
```

```
In [6]: # generators
train_ds=keras.utils.image_dataset_from_directory(
    directory='/content/train',
    labels="inferred",
    label_mode="int",
    batch_size=32,
    image_size=(256, 256)
)

validation_ds=keras.utils.image_dataset_from_directory(
    directory='/content/test',
    labels="inferred",
    label_mode="int",
    batch_size=32,
    image_size=(256, 256)
)
```

Found 20000 files belonging to 2 classes.  
 Found 5000 files belonging to 2 classes.

```
In [7]: # normalized
def process(image,label):
    # cast(x,dtype)
    tf.cast(image/255.,tf.float32)
    return image,label

train_ds=train_ds.map(process)
validation_ds=validation_ds.map(process)
```

```
In [8]: # create CNN model
model=Sequential()
model.add(Conv2D(32,kernel_size=(3,3),padding="valid",activation='relu',input_shape=(256,256,3)))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2),strides=2,padding='valid'))

model.add(Conv2D(64,kernel_size=(3,3),padding="valid",activation='relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2),strides=2,padding="valid"))

model.add(Conv2D(128,kernel_size=(3,3),padding="valid",activation='relu'))
model.add(BatchNormalization())
```

```
model.add(MaxPooling2D(pool_size=(2,2),strides=2,padding="valid"))

model.add(Flatten())

model.add(Dense(128,activation='relu'))
model.add(Dropout(0.1))
model.add(Dense(64,activation='relu'))
model.add(Dropout(0.1))
model.add(Dense(1,activation='sigmoid'))
```

/usr/local/lib/python3.10/dist-packages/keras/src/layers/convolutional/base\_conv.py:107: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

In [9]: model.summary()

**Model: "sequential"**

Layer (type)	Output Shape	Params
conv2d (Conv2D)	(None, 254, 254, 32)	
batch_normalization (BatchNormalization)	(None, 254, 254, 32)	
max_pooling2d (MaxPooling2D)	(None, 127, 127, 32)	
conv2d_1 (Conv2D)	(None, 125, 125, 64)	
batch_normalization_1 (BatchNormalization)	(None, 125, 125, 64)	
max_pooling2d_1 (MaxPooling2D)	(None, 62, 62, 64)	
conv2d_2 (Conv2D)	(None, 60, 60, 128)	
batch_normalization_2 (BatchNormalization)	(None, 60, 60, 128)	
max_pooling2d_2 (MaxPooling2D)	(None, 30, 30, 128)	
flatten (Flatten)	(None, 115200)	
dense (Dense)	(None, 128)	14,848
dropout (Dropout)	(None, 128)	
dense_1 (Dense)	(None, 64)	
dropout_1 (Dropout)	(None, 64)	
dense_2 (Dense)	(None, 1)	

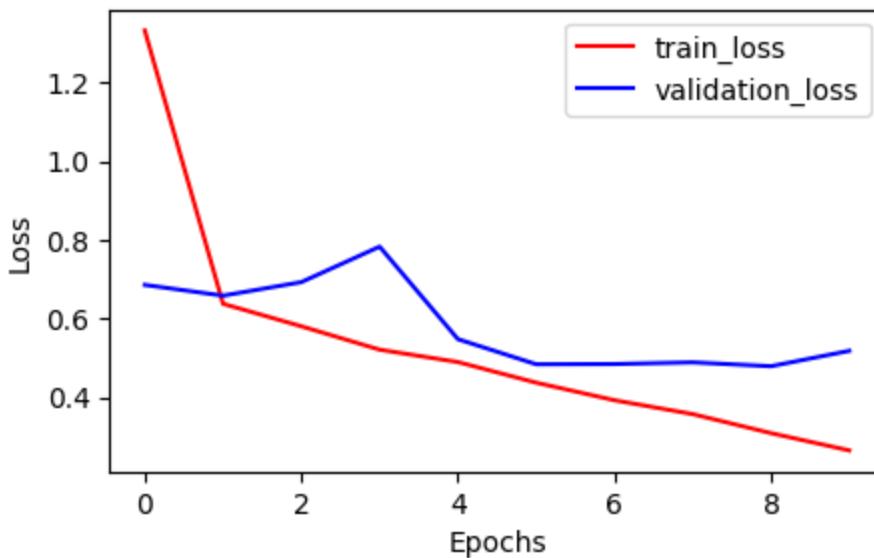
**Total params: 14,848,193 (56.64 MB)**

**Trainable params: 14,847,745 (56.64 MB)**

```
In [10]: model.compile(loss='binary_crossentropy',optimizer='adam',metrics=['accuracy'])
history=model.fit(train_ds,epochs=10,validation_data=validation_ds)
```

```
Epoch 1/10
625/625 65s 84ms/step - accuracy: 0.5539 - loss: 2.3986 - val_accuracy: 0.6114 - val_loss: 0.6851
Epoch 2/10
625/625 50s 80ms/step - accuracy: 0.6287 - loss: 0.6544 - val_accuracy: 0.5918 - val_loss: 0.6578
Epoch 3/10
625/625 81s 78ms/step - accuracy: 0.6998 - loss: 0.5878 - val_accuracy: 0.6598 - val_loss: 0.6921
Epoch 4/10
625/625 83s 80ms/step - accuracy: 0.7489 - loss: 0.5276 - val_accuracy: 0.6592 - val_loss: 0.7825
Epoch 5/10
625/625 81s 78ms/step - accuracy: 0.7690 - loss: 0.5085 - val_accuracy: 0.7354 - val_loss: 0.5471
Epoch 6/10
625/625 83s 80ms/step - accuracy: 0.8033 - loss: 0.4436 - val_accuracy: 0.7830 - val_loss: 0.4832
Epoch 7/10
625/625 85s 84ms/step - accuracy: 0.8323 - loss: 0.3984 - val_accuracy: 0.7908 - val_loss: 0.4838
Epoch 8/10
625/625 82s 85ms/step - accuracy: 0.8528 - loss: 0.3647 - val_accuracy: 0.7794 - val_loss: 0.4881
Epoch 9/10
625/625 82s 84ms/step - accuracy: 0.8781 - loss: 0.3148 - val_accuracy: 0.7992 - val_loss: 0.4783
Epoch 10/10
625/625 82s 84ms/step - accuracy: 0.9006 - loss: 0.2702 - val_accuracy: 0.7950 - val_loss: 0.5175
```

```
In [15]: plt.figure(figsize=(5, 3)) # Specify smaller width and height in inches
plt.plot(history.history['loss'], color='red', label='train_loss')
plt.plot(history.history['val_loss'], color='blue', label='validation_loss')
plt.legend()
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.show()
```



```
In [16]: model=Sequential()
model.add(Conv2D(32,kernel_size=(3,3),padding="valid",activation='relu',input_shape=(28,28,3)))
model.add(MaxPooling2D(pool_size=(2,2),strides=2,padding='valid'))

model.add(Conv2D(64,kernel_size=(3,3),padding="valid",activation='relu',kernel_regularizer=regularizers.l2(0.01)))
model.add(MaxPooling2D(pool_size=(2,2),strides=2,padding="valid"))

model.add(Conv2D(128,kernel_size=(3,3),padding="valid",activation='relu',kernel_regularizer=regularizers.l2(0.01)))
model.add(MaxPooling2D(pool_size=(2,2),strides=2,padding="valid"))

model.add(Flatten())

model.add(Dense(128,activation='relu',kernel_regularizer=tf.keras.regularizers.l2(0.01)))
model.add(Dense(64,activation='relu',kernel_regularizer=tf.keras.regularizers.l2(0.01)))

model.add(Dense(1,activation='sigmoid'))
```

/usr/local/lib/python3.10/dist-packages/keras/src/layers/convolutional/base\_conv.py:107: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

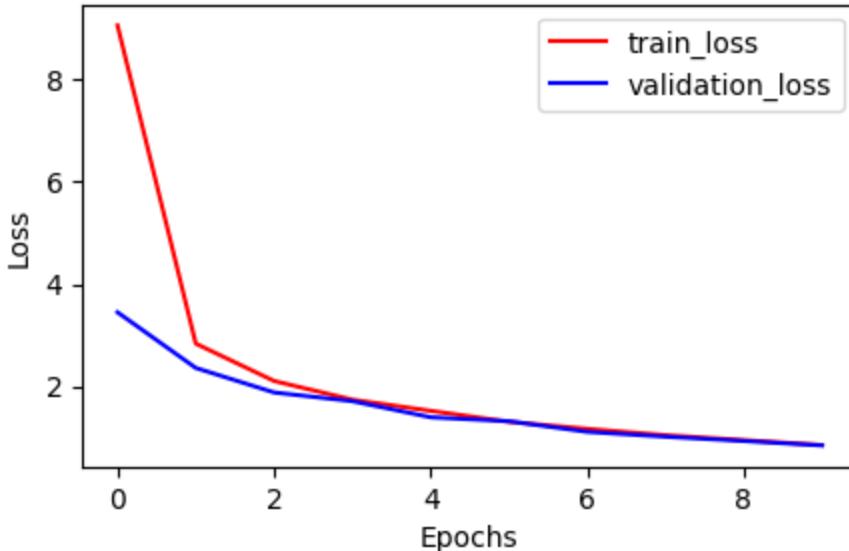
```
In [17]: model.compile(loss='binary_crossentropy',optimizer='Adam',metrics=['accuracy'])
history=model.fit(train_ds,epochs=10,validation_data=validation_ds)
```

```

Epoch 1/10
625/625 52s 76ms/step - accuracy: 0.5469 - loss: 30.1106 - val_accuracy: 0.5786 - val_loss: 3.4499
Epoch 2/10
625/625 78s 74ms/step - accuracy: 0.6065 - loss: 3.1066 - val_accuracy: 0.6422 - val_loss: 2.3612
Epoch 3/10
625/625 81s 73ms/step - accuracy: 0.6469 - loss: 2.2354 - val_accuracy: 0.6552 - val_loss: 1.8818
Epoch 4/10
625/625 82s 73ms/step - accuracy: 0.6677 - loss: 1.8072 - val_accuracy: 0.6216 - val_loss: 1.7119
Epoch 5/10
625/625 82s 74ms/step - accuracy: 0.6823 - loss: 1.6040 - val_accuracy: 0.7328 - val_loss: 1.3967
Epoch 6/10
625/625 89s 85ms/step - accuracy: 0.7377 - loss: 1.3463 - val_accuracy: 0.7160 - val_loss: 1.3196
Epoch 7/10
625/625 80s 82ms/step - accuracy: 0.7590 - loss: 1.2161 - val_accuracy: 0.7730 - val_loss: 1.1123
Epoch 8/10
625/625 79s 78ms/step - accuracy: 0.7804 - loss: 1.0757 - val_accuracy: 0.7772 - val_loss: 1.0171
Epoch 9/10
625/625 78s 71ms/step - accuracy: 0.7923 - loss: 0.9739 - val_accuracy: 0.7792 - val_loss: 0.9337
Epoch 10/10
625/625 83s 73ms/step - accuracy: 0.8051 - loss: 0.8757 - val_accuracy: 0.7824 - val_loss: 0.8463

```

```
In [19]: plt.figure(figsize=(5, 3)) # Adjust to your preferred dimensions
plt.plot(history.history['loss'], color='red', label='train_loss')
plt.plot(history.history['val_loss'], color='blue', label='validation_loss')
plt.legend()
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.show()
```



# RNN

use seattle-weather.csv dataset from Kaggle

```
In [110]: df = pd.read_csv('seattle-weather.csv')
df.head()
```

```
Out[110]:
```

	date	precipitation	temp_max	temp_min	wind	weather
<b>0</b>	2012-01-01	0.0	12.8	5.0	4.7	drizzle
<b>1</b>	2012-01-02	10.9	10.6	2.8	4.5	rain
<b>2</b>	2012-01-03	0.8	11.7	7.2	2.3	rain
<b>3</b>	2012-01-04	20.3	12.2	5.6	4.7	rain
<b>4</b>	2012-01-05	1.3	8.9	2.8	6.1	rain

```
In [111]: df.isnull().sum()
```

```
Out[111]:
```

	<b>0</b>
<b>date</b>	0
<b>precipitation</b>	0
<b>temp_max</b>	0
<b>temp_min</b>	0
<b>wind</b>	0
<b>weather</b>	0

**dtype:** int64

```
In [112]: df.duplicated().sum()
```

```
Out[112]: 0
```

```
In [113]: training_set = df.iloc[:,2:3].values
training_set
```

```
Out[113]: array([[12.8],
 [10.6],
 [11.7],
 ...,
 [ 7.2],
 [ 5.6],
 [ 5.6]])
```

```
In [114]: len(training_set)
```

```
Out[114]: 1461
```

In [115...]

```
def df_to_XY(df,window_size=10):
    X_train=[]
    y_train=[]

    for i in range(10,len(training_set)):
        X_train.append(training_set[i-10:i,0])
        y_train.append(training_set[i,0])

    X_train, y_train = np.array(X_train), np.array(y_train)
    return X_train, y_train
```

In [116...]

```
WINDOW = 10
X,y = df_to_XY(df,WINDOW)
print(len(X),len(y))
X_train = X[:800]
y_train = y[:800]
X_val = X[800:1000]
y_val = y[800:1000]
X_test = X[1000:]
x_test = y[1000:]
```

1451 1451

In [117...]

```
#Reshaping(To add new dimensions)
X_train = np.reshape(X_train,(X_train.shape[0],X_train.shape[1],1))
X_val = np.reshape(X_val,(X_val.shape[0],X_val.shape[1],1))
X_test = np.reshape(X_test,(X_test.shape[0],X_test.shape[1],1))
```

In [118...]

```
#Building the RNN
from keras.models import Sequential
from keras.layers import Dense, LSTM, Dropout
```

In [119...]

```
regressor = Sequential()
```

In [120...]

```
#Addinf the first LSTM Layer and some Dropout regularisation
regressor.add(LSTM(units=50, return_sequences = True, input_shape=(X_train.shape[1], 1))
regressor.add(Dropout(0.2))

regressor.add(LSTM(units=50, return_sequences = True))
regressor.add(Dropout(0.2))

regressor.add(LSTM(units=50, return_sequences = True))
regressor.add(Dropout(0.2))
regressor.add(LSTM(units=50))
regressor.add(Dropout(0.2))
#Output Layer
regressor.add(Dense(units=1))
```

/usr/local/lib/python3.10/dist-packages/keras/src/layers/rnn/rnn.py:204: UserWarning:

Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

In [121...]

```
#Compiling
regressor.compile(optimizer='adam',loss='mean_squared_error')
```

In [130...]

```

from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping
from tensorflow.keras.losses import MeanSquaredError
from tensorflow.keras.metrics import RootMeanSquaredError
from tensorflow.keras.optimizers import Adam

# Fit the model with verbose=0 to suppress output
history = regressor.fit(X_train, y_train, validation_data=(X_val, y_val), epochs=100,

# Display results for the last 5 epochs
last_5_epochs_loss = history.history['loss'][-5:]
last_5_epochs_val_loss = history.history['val_loss'][-5:]

print("Last 5 Epochs' Training and Validation Loss:")
for i, (loss, val_loss) in enumerate(zip(last_5_epochs_loss, last_5_epochs_val_loss),
    epoch_num = 100 - 5 + i # Calculate epoch number for display
    print(f"Epoch {epoch_num:3d}: Training Loss = {loss:.4f}, Validation Loss = {val_1

```

Last 5 Epochs' Training and Validation Loss:

```

Epoch 96: Training Loss = 2.7545, Validation Loss = 13.2436
Epoch 97: Training Loss = 2.4732, Validation Loss = 13.8759
Epoch 98: Training Loss = 2.6289, Validation Loss = 13.3852
Epoch 99: Training Loss = 2.5229, Validation Loss = 13.9444
Epoch 100: Training Loss = 2.8279, Validation Loss = 13.8125

```

In [131...]

```
his = pd.DataFrame(history.history)
```

In [132...]

```
his.head()
```

Out[132]:

	loss	val_loss
0	2.909266	13.461831
1	2.906893	13.124987
2	3.113686	13.059484
3	3.126990	13.089668
4	2.722066	13.163455

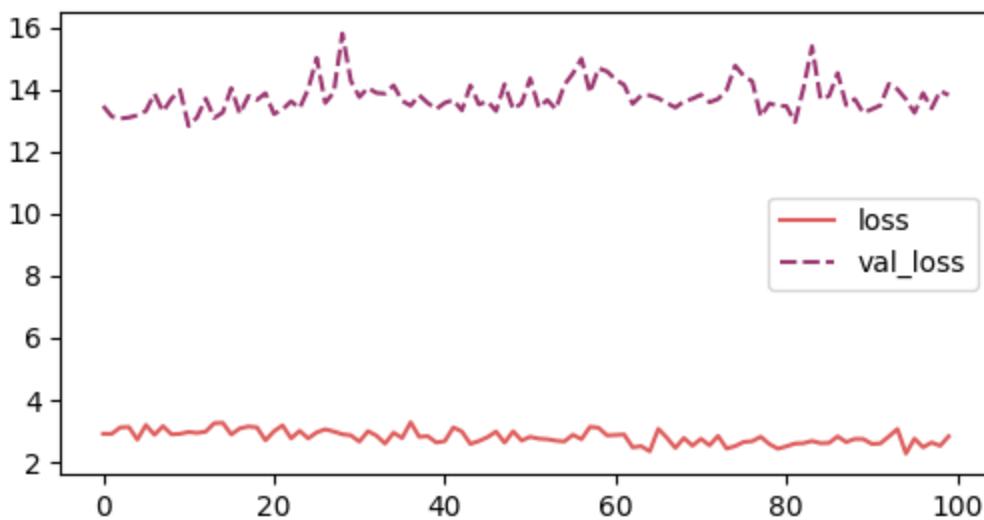
In [138...]

```

import seaborn as sns
his.columns
history_loss = his[['loss', 'val_loss']]
fig,axes = plt.subplots(1,1,figsize=(6,3))
plt.subplot(1,1,1)
plt.title("Loss & Val Loss")
sns.lineplot(history_loss,palette="flare");

```

### Loss & Val Loss



In [139...]

```
train_pred = regressor.predict(X_train).flatten()
val_pred = regressor.predict(X_val).flatten()
test_pred = regressor.predict(X_test).flatten()
```

```
25/25 ━━━━━━━━ 1s 7ms/step
7/7 ━━━━━━━━ 1s 109ms/step
15/15 ━━━━━━ 0s 7ms/step
```

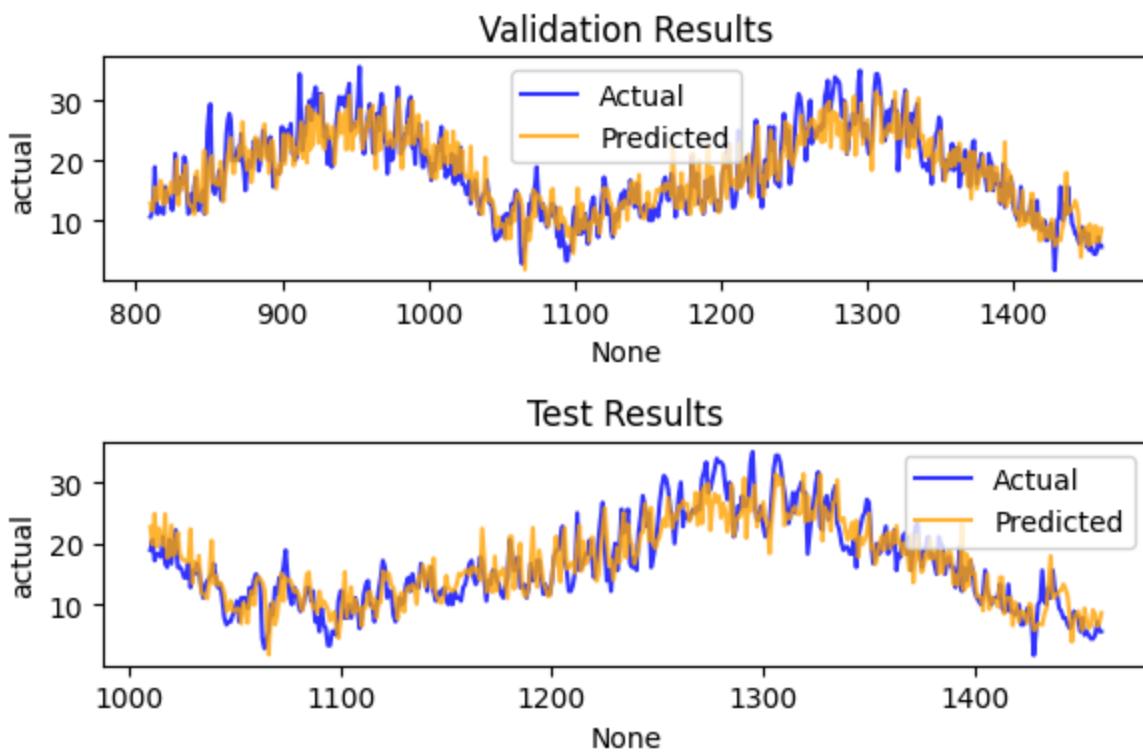
In [143...]

```
fig, axes = plt.subplots(2, 1, figsize=(6, 4), dpi=100)

# Validation Results
plt.subplot(2, 1, 1)
plt.title("Validation Results")
sns.lineplot(data=df_pred[800:], x=df_pred.index[800:], y="actual", color="blue", label="Actual")
sns.lineplot(data=df_pred[800:], x=df_pred.index[800:], y="predicted", color="orange", label="Predicted")

# Test Results
plt.subplot(2, 1, 2)
plt.title("Test Results")
sns.lineplot(data=df_pred[1000:], x=df_pred.index[1000:], y="actual", color="blue", label="Actual")
sns.lineplot(data=df_pred[1000:], x=df_pred.index[1000:], y="predicted", color="orange", label="Predicted")

plt.tight_layout() # Adjust Layout for a better fit
plt.legend() # Show Legend for Labels
plt.show()
```



## LSTM( Long Short Term Memory )

used two dataset from Kaggle 'DailyDelhiClimateTrain.csv' and 'DailyDelhiClimateTest.csv'

```
In [120]: train = pd.read_csv('DailyDelhiClimateTrain.csv')
test = pd.read_csv('DailyDelhiClimateTest.csv')
train.head()
```

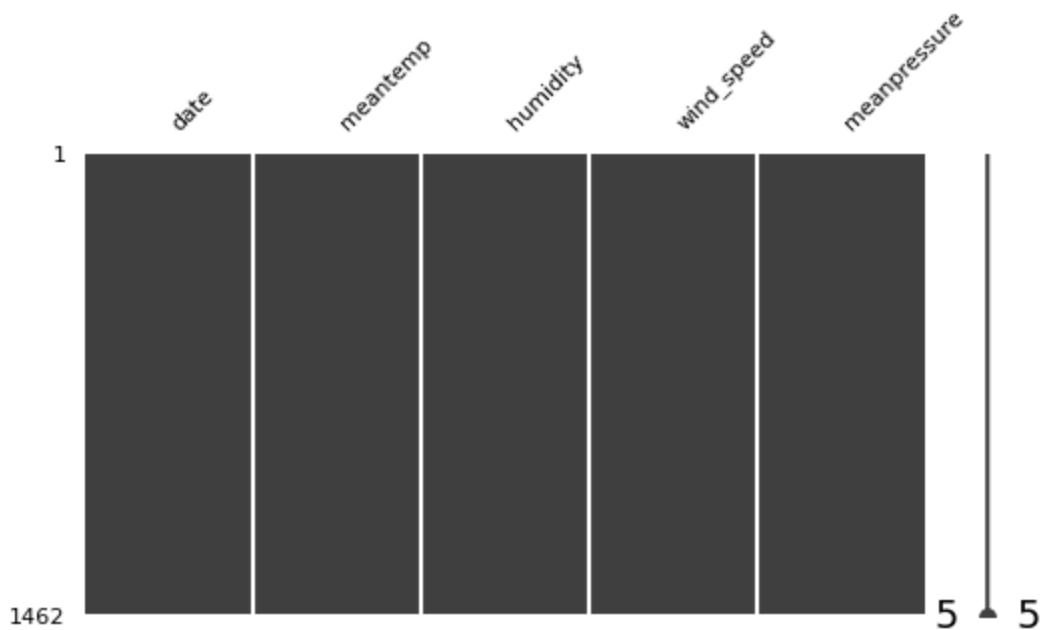
Out[120]:

	date	meantemp	humidity	wind_speed	meanpressure
0	2013-01-01	10.000000	84.500000	0.000000	1015.666667
1	2013-01-02	7.400000	92.000000	2.980000	1017.800000
2	2013-01-03	7.166667	87.000000	4.633333	1018.666667
3	2013-01-04	8.666667	71.333333	1.233333	1017.166667
4	2013-01-05	6.000000	86.833333	3.700000	1016.500000

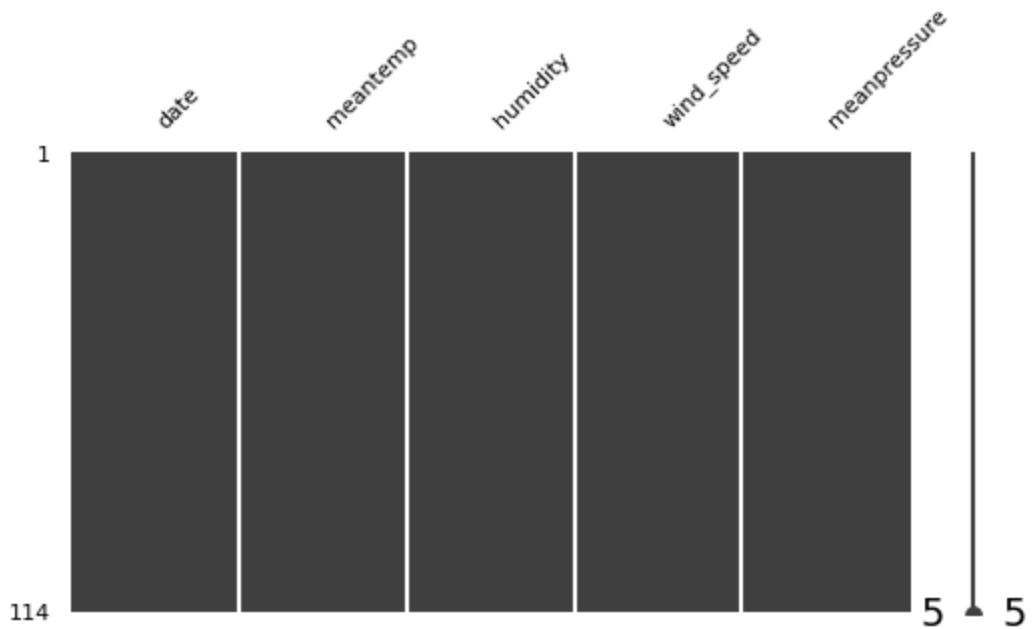
```
In [121]: import missingno as m
plt.figure(figsize=(6, 3))
m.matrix(train, figsize=(6, 3))
plt.tick_params(axis='both', which='major', labelsize=8)
plt.show()

plt.figure(figsize=(6, 3))
m.matrix(test, figsize=(6, 3))
plt.tick_params(axis='both', which='major', labelsize=8)
plt.show()
```

<Figure size 600x300 with 0 Axes>



<Figure size 600x300 with 0 Axes>



In [122...]

```
train.describe()  
test.describe()
```

Out[122]:

	meantemp	humidity	wind_speed	meanpressure
<b>count</b>	114.000000	114.000000	114.000000	114.000000
<b>mean</b>	21.713079	56.258362	8.143924	1004.035090
<b>std</b>	6.360072	19.068083	3.588049	89.474692
<b>min</b>	11.000000	17.750000	1.387500	59.000000
<b>25%</b>	16.437198	39.625000	5.563542	1007.437500
<b>50%</b>	19.875000	57.750000	8.069444	1012.739316
<b>75%</b>	27.705357	71.902778	10.068750	1016.739583
<b>max</b>	34.500000	95.833333	19.314286	1022.809524

In [123...]

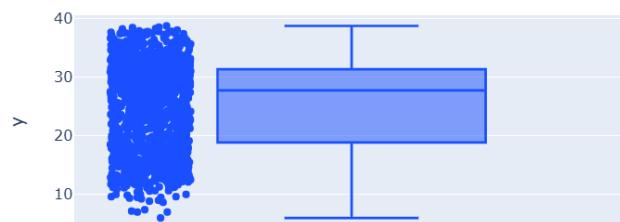
```
print(train.shape)
print(test.shape)
```

(1462, 5)  
(114, 5)

In [134...]

```
#Explanatory Data Analysis (EDA)
import plotly.express as px
fig = px.box(y=train['meantemp'], points='all', title='Temperature', width=600, height=350)
fig.show()
```

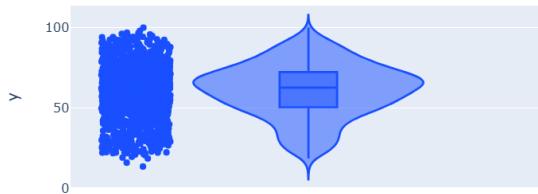
Temperature



In [97]:

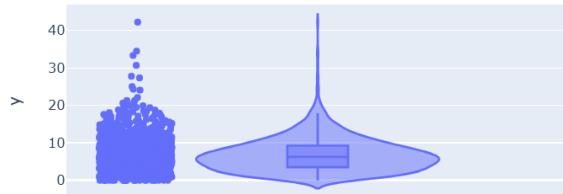
```
fig = px.violin(y=train['humidity'], points='all', box=True, title='Humidity',
                  width=600, height=350,
                  color_discrete_sequence=[ '#194fff' ])
fig.show()
```

Humidity



```
In [98]: fig = px.violin(y=train['wind_speed'], points='all', box=True, title='Wind Speed', width=600)
fig.show()
```

Wind Speed



```
In [99]: #Temperature
import plotly.graph_objects as go
fig = go.Figure()
fig.add_trace(go.Scatter(x=train['date'], y=train['meantemp'], name='Real Price',
                         line=dict(color='royalblue', width=2)))

fig.update_layout(title='Temperature', width=600, height=350) # Adjust width and height
fig.show()
```

Temperature



In [100]:

```
import plotly.graph_objects as go
fig2 = go.Figure()

fig2.add_trace(go.Scatter( x=train['date'], y=train['humidity'], name = 'Predicted Pri
                           line=dict(color='firebrick', width=2)))
fig.update_layout(title='Humidity', width=600, height=350)
fig.show()
```

Humidity



In [79]:

```
temp_train = train.iloc[:,1:2]
temp_test = test.iloc[:,1:2]
```

```
In [80]: from sklearn.preprocessing import MinMaxScaler
ss= MinMaxScaler(feature_range=(0,1))
temp_train= ss.fit_transform(temp_train)
temp_test= ss.fit_transform(temp_test)
```

```
In [81]: def create_dataset(dataset, look_back=1):
    dataX, dataY = [], []
    for i in range(len(dataset)-look_back-1):
        a = dataset[i:(i+look_back), 0]
        dataX.append(a)
        dataY.append(dataset[i + look_back, 0])
    return np.array(dataX), np.array(dataY)
```

```
In [82]: look_back = 1
trainX, trainY = create_dataset(temp_train, look_back)
testX, testY = create_dataset(temp_test, look_back)
```

```
In [83]: trainX = np.reshape(trainX, (trainX.shape[0], trainX.shape[1],1))
testX = np.reshape(testX, (testX.shape[0], testX.shape[1],1))
```

### Building the LSTM model

```
In [84]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Input, LSTM, Dropout, Dense

# Initialize the model
model_temp = Sequential()
model_temp.add(Input(shape=(trainX.shape[1], 1)))
model_temp.add(LSTM(units=100, return_sequences=True))
model_temp.add(Dropout(0.2))

model_temp.add(LSTM(units=100, return_sequences=True))
model_temp.add(Dropout(0.2))

model_temp.add(LSTM(units=100, return_sequences=True))
model_temp.add(Dropout(0.2))

model_temp.add(LSTM(units=50))
model_temp.add(Dropout(0.2))
model_temp.add(Dense(units=1))
model_temp.compile(optimizer='adam', loss='mean_squared_error')
```

```
In [85]: #Compiling with ADAM Optimizer and Mean Squared Error
# Compiling the RNN
model_temp.compile(optimizer = 'adam', loss = 'mean_squared_error', metrics=[ 'accuracy']
# Fitting the RNN to the Training set
history = model_temp.fit(trainX, trainY, epochs = 100, batch_size = 32,verbose=0)
for i in range(95, 100):
    print(f"Epoch {i+1}/100 - loss: {history.history['loss'][i]:.4f} - accuracy: {hist
```

Epoch 96/100 - loss: 0.0034 - accuracy: 0.0014  
 Epoch 97/100 - loss: 0.0036 - accuracy: 0.0014  
 Epoch 98/100 - loss: 0.0034 - accuracy: 0.0014  
 Epoch 99/100 - loss: 0.0035 - accuracy: 0.0014  
 Epoch 100/100 - loss: 0.0036 - accuracy: 0.0014

```
In [86]: #Prediction for Test Set
```

```
prediction = model_temp.predict(testX)
prediction = ss.inverse_transform(prediction)
temp_test = ss.inverse_transform(temp_test)
#prediction
```

```
4/4 ━━━━━━━━ 1s 193ms/step
```

```
In [87]: temp_test.shape
```

```
Out[87]: (114, 1)
```

```
In [88]: prediction.shape
```

```
Out[88]: (112, 1)
```

```
In [89]: prediction1 = prediction
```

```
In [90]: prediction.mean()
```

```
Out[90]: 21.748611
```

```
In [91]: prediction1 = np.append(prediction1, [21.90234,21.90234])
```

```
In [92]: prediction1
```

```
prediction1 = prediction1.reshape(len(prediction1),1)
prediction1.shape
```

```
Out[92]: (114, 1)
```

```
In [93]: data = np.concatenate((prediction1.reshape(len(prediction1),1), temp_test.reshape(len(temp_test),1)))
#print(data)
```

```
In [94]: from sklearn.metrics import r2_score
meantemp = r2_score(temp_test, prediction1)
print("Accuracy of Temperature : ", meantemp)
```

```
Accuracy of Temperature : 0.9395911399801052
```